

Control variates for a GARCH model

A Monte Carlo Study with Zero-Variance Estimators

Ralph Nader Kenan Darfilali Tristan Picard-Bernes

ENSAE Paris
Simulation & Monte Carlo Methods

May 2, 2026

1. Introduction
2. Model and Bayesian Setup
3. Sampling the Posterior
4. Variance Reduction: Zero-Variance Control Variates
5. Bonus: Validity of OLS on a Correlated Chain
6. Conclusion

PART 1

Introduction

Motivation

The core problem

Bayesian inference on GARCH models requires evaluating intractable posterior expectations, so we rely on **Markov Chain Monte Carlo (MCMC)**. But MCMC estimators can have **high variance**, requiring very long chains to be reliable.

- **Goal:** estimate the posterior mean of $\theta = (\omega_1, \omega_2, \omega_3)$ in a Normal-GARCH(1,1) model efficiently.
- **Key idea:** replace f by a corrected estimator $\tilde{f} = f + \mathbf{a}^T \mathbf{z}$ with the same expectation under π but much smaller variance, using **zero-variance control variates**.
- **Extension:** as the number of control variates grows, we use **LASSO** for dimensionality, and **subsampling** / **block averaging** to validate the regression step.

Based on Mira, Solgi & Imparato, *Stat. Comput.* 23 (2013), 653–662, applied to the DEM/GBP exchange rate data of Ardia (2008).

PART 2

Model and Bayesian Setup

The Model — Normal-GARCH(1,1)

Definition

The return process $\{r_t\}$ follows

$$r_t \mid \mathcal{F}_{t-1} \sim \mathcal{N}(0, h_t), \quad h_t = \omega_1 + \omega_2 r_{t-1}^2 + \omega_3 h_{t-1}$$

with $\omega_1 > 0$, $\omega_2, \omega_3 \geq 0$, $\omega_2 + \omega_3 < 1$ (covariance stationarity).

- ω_1 — long-run variance floor (background volatility)
- ω_2 — ARCH term: reaction to yesterday's shock
- ω_3 — GARCH term: variance persistence (memory)
- Initialised at $h_0 = \omega_1 / (1 - \omega_2 - \omega_3)$

Bayesian Setup

Prior


Independent truncated Gaussians: $\omega_j \sim \mathcal{N}(0, 1) \mathbf{1}_{\omega_j \geq 0}$, $j = 1, 2, 3$. Weakly informative — posterior is driven by the data.

Posterior

By Bayes' theorem:

$$\pi(\theta \mid r_{1:T}) \propto \underbrace{\exp\left(-\frac{1}{2} \sum_{t=1}^T \left[\log h_t + \frac{r_t^2}{h_t}\right]\right)}_{\text{likelihood}} \times \underbrace{\pi(\theta)}_{\text{prior}}$$

The Simulated Data



$T = 3000$ returns simulated with $\theta = (0.055, 0.245, 0.600)$.

- Returns oscillate around zero (zero-mean by construction) with occasional spikes exceeding $\pm 2.5\%$.
- Visible **volatility clustering** — alternating calm and turbulent stretches, the GARCH signature.
- Conditional variance h_t never reaches zero (floor ~ 0.2 from ω_1); spikes in h_t align with bursts in r_t .
- Unconditional std: $\sqrt{0.055/(1 - 0.245 - 0.600)} \approx 0.6\%$ per day.

PART 3

Sampling the Posterior

The Sampler — Random Walk Metropolis

Why RWM?


The posterior $\pi(\theta \mid r) \propto L(\theta \mid r) \pi(\theta)$ has **no closed form** since the likelihood involves the recursive variance h_t . RWM only needs to *evaluate* π at two points per iteration.

- **Proposal:** $\theta' = \theta + \varepsilon$, $\varepsilon \sim \mathcal{N}(\mathbf{0}, \tau^2 I_3)$
- **Acceptance:** $\min(1, \pi(\theta' \mid r) / \pi(\theta \mid r))$
- Computed in log-scale to avoid underflow
- Positivity enforced via $\log \pi = -\infty$ outside the support

Calibrating τ

Aim for acceptance rate $\approx$ **25%**: too high means the chain barely moves whereas too low means proposals are rejected and the chain freezes.

RWM on Simulated Data




- $T = 3000$ simulated returns,
 $\theta = (0.055, 0.245, 0.600)$
- $\tau = 0.01$, acceptance = **27.1%** ✓
- $N = 12,000$, burn-in = 2,000
- **Identification:** ω_1 & ω_2 partially compensate; ω_3 tightest.

	ω_1	ω_2	ω_3
True	0.055	0.245	0.600
Posterior	0.0585	0.2285	0.5972

Posterior recovery.

The Dependence Problem



Autocorrelation decays slowly: ACF still ≈ 0.4 at lag 100.

Effective Sample Size


Out of 10,000 post-burn-in draws:

- ESS $\omega_1 = 70$
- ESS $\omega_2 = 76$
- ESS $\omega_3 = 53$

~ 1 independent sample / 140 iterations.

- Root cause: isotropic proposal does not adapt to the posterior ridge.
- MCMC variance $\gg$ iid case $\Rightarrow$ control variates are needed.

Real Data — DEM/GBP 1985–1987




756 daily returns of the DEM/GBP exchange rate (Ardia, 2008).

- Same dataset used by Mira et al. (2013) for direct comparison.
- Visible volatility clustering: calm stretches alternating with turbulent ones — the empirical motivation for the GARCH model.
- Mean return $\approx 0\%$, std $\approx 0.57\%$ per day.

RWM on Real Data

RWM on DEM/GBP from 1985–1987



Trace plots and posterior histograms on 756 daily DEM/GBP returns.

- $\tau = 0.02$, acceptance = **23.1%** ✓
- $N = 50,000$, burn-in = 5,000
- $\hat{\omega}_2 + \hat{\omega}_3 \approx 0.84$: stationary but persistent volatility

	ω_1	ω_2	ω_3
Ours	0.0545	0.2335	0.6027
Paper	0.055	0.245	0.600

Implementation validated.

PART 4

Variance Reduction: Zero-Variance Control Variates

The Zero-Variance Principle (Mira et al. 2013)

Idea

Replace $f(\theta)$ with a corrected estimator $\tilde{f}(\theta) = f(\theta) + \mathbf{a}^T \mathbf{z}(\theta)$, where $\mathbf{z}$ is **mean-zero** under π . Then $\mathbb{E}_\pi[\tilde{f}] = \mathbb{E}_\pi[f]$: same target but lower variance if $\mathbf{a}$ is chosen well.

Natural choice of $\mathbf{z}$ — the score: $\mathbf{z}(\theta) = -\frac{1}{2} \nabla \log \pi(\theta \mid r)$, $\mathbb{E}_\pi[\mathbf{z}] = \mathbf{0}$.


Optimal $\mathbf{a}$ — found by **OLS regression** of f on $\mathbf{z}$ along the chain:

$$f(\theta^i) = \alpha + \mathbf{a}^T \mathbf{z}(\theta^i) + \varepsilon_i$$

Intercept $\hat{\alpha}$ is the ZV estimate of $\mathbb{E}_\pi[f]$; residual variance gives $\text{Var}(\tilde{f})$.

For GARCH: $\mathbf{z}$ computed by differentiating the variance recursion $\dot{h}_t^{(j)} = \partial h_t / \partial \omega_j$ via the chain rule.

Zero-Variance Degree-1 — Simulated Data




Posteriors visibly tighten: the ZV correction collapses the bulk of the distribution.

	MC std error (naive $\rightarrow$ ZV)	Reduction
ω_1	0.009 $\rightarrow$ 0.003	12.0×
ω_2	0.025 $\rightarrow$ 0.005	22.2×
ω_3	0.040 $\rightarrow$ 0.010	15.4×

Std errors in units of 10^{-2} . All within the 95% CI of Mira et al. (2013).

Zero-Variance Degree-1 — Real Data (DEM/GBP)



Same pattern, smaller magnitude: the linear $\mathbf{a}^T \mathbf{z}$ explains less of the irregular real-data posterior.

	MC std error (naive $\rightarrow$ ZV)	Reduction
ω_1	0.007 $\rightarrow$ 0.004	3.6×
ω_2	0.023 $\rightarrow$ 0.009	6.1×
ω_3	0.034 $\rightarrow$ 0.016	4.8×

Real data underperforms $\Rightarrow$ motivates degree-2.

Scaling Up — Degree-2 Polynomials

The 9 control variates

With $P(\theta) = \mathbf{a}^T \theta + \frac{1}{2} \theta^T B \theta$:

$$\tilde{f}(\theta) = f(\theta) - \frac{1}{2} \text{tr}(B) + (\mathbf{a} + B\theta)^T \mathbf{z}(\theta)$$

- 3 linear: z_1, z_2, z_3
- 6 quadratic: $\omega_j z_k$ for $j \leq k$

Why not just OLS?

Cost $\mathcal{O}(nm^2 + m^3)$, $m = \frac{d(d+3)}{2}$:

Degree	$d = 3$	$d = 10$	$d = 50$
1	3	10	50
2	9	65	1325
3	19	285	23,425


Prohibitive at moderate d .

LASSO (Leluc et al. 2021)

ℓ_1 penalty: $\min_{\mathbf{a}} \|f - H\mathbf{a}\|^2 + \lambda \|\mathbf{a}\|_1$

- Selects relevant variates
- λ tuned by 5-fold CV
- Avoids overfitting on autocorrelated draws
- Robust at large m

Degree-2 LASSO — Simulated Data




Degree-2 LASSO (green) posteriors collapse onto the true values (dashed) — invisibly thin in the histogram.

Method	ω_1	ω_2	ω_3
Naive	1×	1×	1×
Deg-1 OLS	12×	22×	15×
Deg-2 LASSO	738×	422×	1028×

LASSO selects 9/9 (ω_1, ω_3) and 8/9 (ω_2) variates — no sparsity at $d = 3$.

Degree-2 LASSO — Real Data (DEM/GBP)



Same dramatic collapse on real financial data.

Method	ω_1	ω_2	ω_3
Naive	1×	1×	1×
Deg-1 OLS	3.6×	6.1×	4.8×
Deg-2 LASSO	298×	340×	760×

Three orders of magnitude reduction: it confirms the paper's central finding.

PART 5


Bonus: Validity of OLS on a Correlated Chain

Bonus — Is OLS Valid on a Correlated Chain?


The issue

OLS assumes **i.i.d.** observations. With autocorrelated MCMC draws: $\hat{\Sigma}_{zz}$ uses fewer effective samples than N , coefficients $\hat{a}$ are less precise, and LASSO cross-validation underestimates uncertainty (adjacent folds nearly identical).

- **Sub-sampling**: keep every n/ESS -th draw
- **Block averaging**: replace each block by its mean



Simulated data — both methods help, block avg triples the reduction.



Real data — sub-sampling barely helps, only block averaging delivers.

Conclusion: block averaging is the robust decorrelation strategy. Moreover, the Q2/Q3 results were *conservative*, not optimistic.

PART 6

Conclusion

Conclusion

- Implemented Bayesian Normal-GARCH(1,1) estimation via RWM with posterior means that do match Mira et al. (2013) within 1%.
- **Degree-1 ZV control variates:** 12–22 $\times$ variance reduction (simulated), 4–6 $\times$ (real).
- **Degree-2 + LASSO: three orders of magnitude** reduction ($\sim 10^3\times$), confirming the paper's finding and extending it via Leluc et al. (2021).
- Addressed OLS validity on correlated MCMC: block averaging shows our reductions were **conservative**.

Possible extensions

Higher-degree polynomials, adaptive MCMC for better mixing, multivariate or asymmetric GARCH (E-GARCH, GJR-GARCH).

Thank you.

References I

- [1] Mira, A., Solgi, R., & Imparato, D. (2013). *Zero variance Markov chain Monte Carlo for Bayesian estimators*. *Statistics and Computing*, 23(5), 653–662.
- [2] Leluc, R., Portier, F., & Segers, J. (2021). *Control variate selection for Monte Carlo integration*. *Statistics and Computing*, 31(4), 1–26.
- [3] Bollerslev, T. (1986). *Generalized autoregressive conditional heteroskedasticity*. *Journal of Econometrics*, 31(3), 307–327.
- [4] Ardia, D. (2008). *Financial Risk Management with Bayesian Estimation of GARCH Models*. Springer.